

# Inference and Representation

DS-GA-1005, CSCI-GA.2569

David Sontag

New York University

Lecture 2, Sept. 19, 2016

# Today's lecture

- Markov random fields
  - ① Formalism
  - ② Conditional independence properties
  - ③ Examples
- Factor graphs
- Bayesian networks  $\Rightarrow$  Markov random fields (*moralization*)
- Learning graphical models from data

# Bayesian networks

## Reminder of first lecture

- A **Bayesian network** is specified by a directed *acyclic* graph  $G = (V, E)$  with:
  - 1 One node  $i \in V$  for each random variable  $X_i$
  - 2 One conditional probability distribution (CPD) per node,  $p(x_i \mid \mathbf{x}_{\text{Pa}(i)})$ , specifying the variable's probability conditioned on its parents' values
- Corresponds 1-1 with a particular factorization of the joint distribution:

$$p(x_1, \dots, x_n) = \prod_{i \in V} p(x_i \mid \mathbf{x}_{\text{Pa}(i)})$$


- Powerful framework for designing *algorithms* to perform probability computations

# Bayesian networks have limitations

- Recall that  $G$  is a **perfect map** for distribution  $p$  if  $I(G) = I(p)$
- Theorem:** Not every distribution has a perfect map as a DAG

## Proof.

(By counterexample.) There is a distribution on 4 variables where the only independencies are  $A \perp C \mid \{B, D\}$  and  $B \perp D \mid \{A, C\}$ . This cannot be represented by any Bayesian network.



Both (a) and (b) encode  $(A \perp C \mid B, D)$ , but in both cases  $(B \not\perp D \mid A, C)$ . □

# Example

- Let's come up with an example of a distribution  $p$  satisfying  $A \perp C \mid \{B, D\}$  and  $B \perp D \mid \{A, C\}$
- $A$ =Alex's hair color (red, green, blue)  
 $B$ =Bob's hair color  
 $C$ =Catherine's hair color  
 $D$ =David's hair color
- Alex and Bob are friends, Bob and Catherine are friends, Catherine and David are friends, David and Alex are friends
- Friends *never* have the same hair color!

# Undirected graphical models

- An alternative representation for joint distributions is as an **undirected graphical model**
- As in BNs, we have one node for each random variable
- Rather than CPDs, we specify (non-negative) **potential functions** over sets of variables associated with cliques  $C$  of the graph,

$$p(x_1, \dots, x_n) = \frac{1}{Z} \prod_{c \in C} \phi_c(\mathbf{x}_c)$$

$Z$  is the **partition function** and normalizes the distribution:

$$Z = \sum_{\hat{x}_1, \dots, \hat{x}_n} \prod_{c \in C} \phi_c(\hat{\mathbf{x}}_c)$$

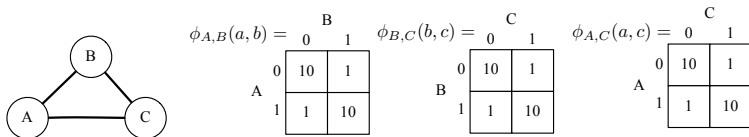
- Like CPD's,  $\phi_c(\mathbf{x}_c)$  can be represented as a table, but it is *not normalized*
- Also known as **Markov random fields** (MRFs) or Markov networks

# Undirected graphical models

$$p(x_1, \dots, x_n) = \frac{1}{Z} \prod_{c \in C} \phi_c(\mathbf{x}_c),$$

$$Z = \sum_{\hat{x}_1, \dots, \hat{x}_n} \prod_{c \in C} \phi_c(\hat{\mathbf{x}}_c)$$

Simple example (potential function on each edge encourages the variables to take the same value):



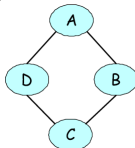
$$p(a, b, c) = \frac{1}{Z} \phi_{A,B}(a, b) \cdot \phi_{B,C}(b, c) \cdot \phi_{A,C}(a, c),$$

where

$$Z = \sum_{\hat{a}, \hat{b}, \hat{c} \in \{0,1\}^3} \phi_{A,B}(\hat{a}, \hat{b}) \cdot \phi_{B,C}(\hat{b}, \hat{c}) \cdot \phi_{A,C}(\hat{a}, \hat{c}) = 2 \cdot 1000 + 6 \cdot 10 = 2060.$$

# Hair color example as a MRF

- We now have an **undirected** graph:



- The joint probability distribution is parameterized as

$$p(a, b, c, d) = \frac{1}{Z} \phi_{AB}(a, b) \phi_{BC}(b, c) \phi_{CD}(c, d) \phi_{AD}(a, d) \phi_A(a) \phi_B(b) \phi_C(c) \phi_D(d)$$

- **Pairwise potentials** enforce that no friend has the same hair color:

$$\phi_{AB}(a, b) = 0 \text{ if } a = b, \text{ and } 1 \text{ otherwise}$$

- **Single-node potentials** specify an affinity for a particular hair color, e.g.

$$\phi_D(\text{"red"}) = 0.6, \quad \phi_D(\text{"blue"}) = 0.3, \quad \phi_D(\text{"green"}) = 0.1$$

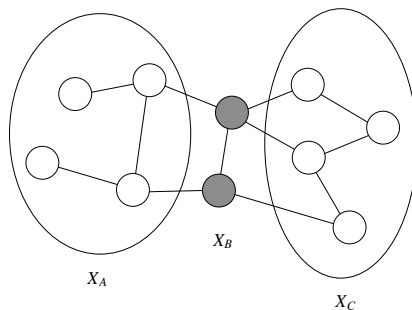
The normalization  $Z$  makes the potentials **scale invariant**! Equivalent to

$$\phi_D(\text{"red"}) = 6, \quad \phi_D(\text{"blue"}) = 3, \quad \phi_D(\text{"green"}) = 1$$



# Markov network structure implies conditional independencies


- Let  $G$  be the undirected graph where we have one edge for every pair of variables that appear together in a potential
- Conditional independence is given by **graph separation!**



- $X_A \perp X_C \mid X_B$  if there is no path from  $a \in \mathbf{A}$  to  $c \in \mathbf{C}$  after removing all variables in  $\mathbf{B}$

# Example

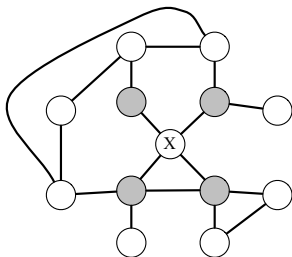
- Returning to hair color example, its undirected graphical model is:



- Since removing  $A$  and  $C$  leaves no path from  $D$  to  $B$ , we have  $D \perp B \mid \{A, C\}$
- Similarly, since removing  $D$  and  $B$  leaves no path from  $A$  to  $C$ , we have  $A \perp C \mid \{D, B\}$
- No other independencies implied by the graph

# Markov blanket


- A set  $\mathbf{U}$  is a **Markov blanket** of  $X$  if  $X \notin \mathbf{U}$  and if  $\mathbf{U}$  is a minimal set of nodes such that  $X \perp (\mathcal{X} - \{X\} - \mathbf{U}) \mid \mathbf{U}$
- In undirected graphical models, the Markov blanket of a variable is precisely its **neighbors** in the graph:



- In other words,  $X$  is independent of the rest of the nodes in the graph given its immediate neighbors

# Proof of independence through separation

- We will show that  $A \perp C \mid B$  for the following distribution:


$$p(a, b, c) = \frac{1}{Z} \phi_{AB}(a, b) \phi_{BC}(b, c)$$


- First, we show that  $p(a \mid b)$  can be computed using only  $\phi_{AB}(a, b)$ :

$$\begin{aligned} p(a \mid b) &= \frac{p(a, b)}{p(b)} \\ &= \frac{\frac{1}{Z} \sum_{\hat{c}} \phi_{AB}(a, b) \phi_{BC}(b, \hat{c})}{\frac{1}{Z} \sum_{\hat{a}, \hat{c}} \phi_{AB}(\hat{a}, b) \phi_{BC}(b, \hat{c})} \\ &= \frac{\phi_{AB}(a, b) \sum_{\hat{c}} \phi_{BC}(b, \hat{c})}{\sum_{\hat{a}} \phi_{AB}(\hat{a}, b) \sum_{\hat{c}} \phi_{BC}(b, \hat{c})} = \frac{\phi_{AB}(a, b)}{\sum_{\hat{a}} \phi_{AB}(\hat{a}, b)}. \end{aligned}$$

- More generally, the probability of a variable conditioned on its Markov blanket depends *only* on potentials involving that node

# Proof of independence through separation

- We will show that  $A \perp C \mid B$  for the following distribution:



$$p(a, b, c) = \frac{1}{Z} \phi_{AB}(a, b) \phi_{BC}(b, c)$$


Proof.

$$\begin{aligned} p(a, c \mid b) &= \frac{p(a, c, b)}{\sum_{\hat{a}, \hat{c}} p(\hat{a}, b, \hat{c})} = \frac{\phi_{AB}(a, b) \phi_{BC}(b, c)}{\sum_{\hat{a}, \hat{c}} \phi_{AB}(\hat{a}, b) \phi_{BC}(b, \hat{c})} \\ &= \frac{\phi_{AB}(a, b) \phi_{BC}(b, c)}{\sum_{\hat{a}} \phi_{AB}(\hat{a}, b) \sum_{\hat{c}} \phi_{BC}(b, \hat{c})} \\ &= p(a \mid b) p(c \mid b) \end{aligned}$$



# Example: Ising model


- Invented by the physicist Wilhelm Lenz (1920), who gave it as a problem to his student Ernst Ising
- Mathematical model of ferromagnetism in statistical mechanics
- The spin of an atom is biased by the spins of atoms nearby on the material:



- Each atom  $X_i \in \{-1, +1\}$ , whose value is the direction of the atom spin
- If a spin at position  $i$  is  $+1$ , what is the probability that the spin at position  $j$  is also  $+1$ ?
- Are there phase transitions where spins go from “disorder” to “order”?

## Example: Ising model

- Each atom  $X_i \in \{-1, +1\}$ , whose value is the direction of the atom spin
- The spin of an atom is biased by the spins of atoms nearby on the material:



$$p(x_1, \dots, x_n) = \frac{1}{Z} \exp \left( \sum_{i < j} w_{i,j} x_i x_j - \sum_i u_i x_i \right)$$

- When  $w_{i,j} > 0$ , nearby atoms encouraged to have the same spin (called **ferromagnetic**), whereas  $w_{i,j} < 0$  encourages  $X_i \neq X_j$
- Node potentials  $\exp(-u_i x_i)$  encode the bias of the individual atoms
- Scaling the parameters makes the distribution more or less spiky

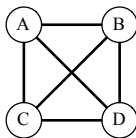
# Higher-order potentials

- The examples so far have all been **pairwise MRFs**, involving only node potentials  $\phi_i(X_i)$  and pairwise potentials  $\phi_{i,j}(X_i, X_j)$
- Often we need **higher-order** potentials, e.g.

$$\phi(x, y, z) = 1[x + y + z \geq 1],$$

where  $X, Y, Z$  are binary, enforcing that at least one of the variables takes the value 1

- Although Markov networks are useful for understanding independencies, they hide much of the distribution's structure:

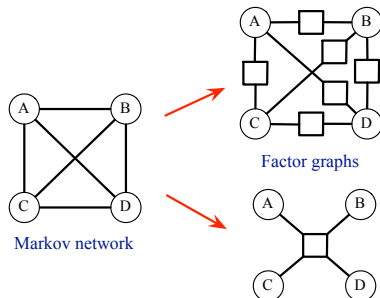


Does this have pairwise potentials, or one potential for all 4 variables?



# Factor graphs

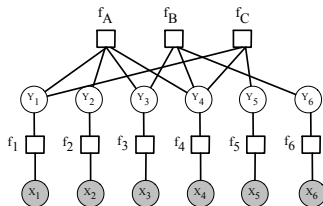
- $G$  does not reveal the structure of the distribution: maximum cliques vs. subsets of them
- A **factor graph** is a bipartite undirected graph with variable nodes and factor nodes. Edges are only between the variable nodes and the factor nodes
- Each factor node is associated with a single potential, whose scope is the set of variables that are neighbors in the factor graph



- The distribution is same as the MRF – this is just a different data structure

# Example: Low-density parity-check codes

- Error correcting codes for transmitting a message over a noisy channel (invented by Gallager in the 1960's, then re-discovered in 1996)



- Each of the top row factors enforce that its variables have even parity:


$$f_A(Y_1, Y_2, Y_3, Y_4) = 1 \text{ if } Y_1 \otimes Y_2 \otimes Y_3 \otimes Y_4 = 0, \text{ and } 0 \text{ otherwise}$$

- Thus, the only assignments  $\mathbf{Y}$  with non-zero probability are the following (called **codewords**): *3 bits encoded using 6 bits*

000000, 011001, 110010, 101011, 111100, 100101, 001110, 010111

- $f_i(Y_i, X_i) = p(X_i | Y_i)$ , the likelihood of a bit flip according to noise model

# Example: Low-density parity-check codes



- The *decoding* problem for LDPCs is to find

$$\operatorname{argmax}_{\mathbf{y}} p(\mathbf{y} \mid \mathbf{x})$$

This is called the **maximum a posteriori** (MAP) assignment

- Since  $Z$  and  $p(\mathbf{x})$  are constants with respect to the choice of  $\mathbf{y}$ , can equivalently solve (taking the log of  $p(\mathbf{y}, \mathbf{x})$ ):

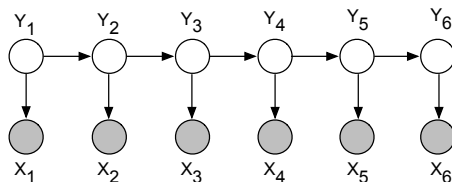
$$\operatorname{argmax}_{\mathbf{y}} \sum_{c \in C} \theta_c(\mathbf{y}_c, \mathbf{x}_c),$$

where  $\theta_c(\mathbf{x}_c) = \log \phi_c(\mathbf{y}_c, \mathbf{x}_c)$

- This is a discrete optimization problem!

# Converting BNs to Markov networks

What is the equivalent Markov network for a hidden Markov model?



Many inference algorithms are more conveniently given for undirected models – this shows how they can be applied to Bayesian networks

# Moralization of Bayesian networks

- Procedure for converting a Bayesian network into a Markov network
- The **moral graph**  $\mathcal{M}[G]$  of a BN  $G = (V, E)$  is an undirected graph over  $V$  that contains an undirected edge between  $X_i$  and  $X_j$  if
  - 1 there is a directed edge between them (in either direction)
  - 2  $X_i$  and  $X_j$  are both parents of the same node



(term historically arose from the idea of “marrying the parents” of the node)

- The addition of the moralizing edges leads to the loss of some independence information, e.g.,  $A \rightarrow C \leftarrow B$ , where  $A \perp B$  is lost

# Converting BNs to Markov networks

- 1 Moralize the directed graph to obtain the undirected graphical model:



- 2 Introduce one potential function for each CPD:

$$\phi_i(x_i, \mathbf{x}_{pa(i)}) = p(x_i \mid \mathbf{x}_{pa(i)})$$

- So, converting a hidden Markov model to a Markov network is simple:



- For variables having  $> 1$  parent, factor graph notation is useful

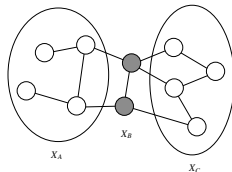
# Factorization implies conditional independencies

- $p(\mathbf{x})$  is a *Gibbs distribution* over  $G$  if it can be written as

$$p(x_1, \dots, x_n) = \frac{1}{Z} \prod_{c \in C} \phi_c(\mathbf{x}_c),$$

where the variables in each potential  $c \in C$  form a clique in  $G$

- Recall that conditional independence is given by graph separation:



- Theorem (**soundness of separation**): If  $p(\mathbf{x})$  is a Gibbs distribution for  $G$ , then  $G$  is an I-map for  $p(\mathbf{x})$ , i.e.  $I(G) \subseteq I(p)$

*Proof:* Suppose  $\mathbf{B}$  separates  $\mathbf{A}$  from  $\mathbf{C}$ . Then we can write

$$p(\mathbf{X}_A, \mathbf{X}_B, \mathbf{X}_C) = \frac{1}{Z} f(\mathbf{X}_A, \mathbf{X}_B) g(\mathbf{X}_B, \mathbf{X}_C).$$

# Conditional independencies implies factorization

- Theorem (**soundness of separation**): If  $p(\mathbf{x})$  is a Gibbs distribution for  $G$ , then  $G$  is an I-map for  $p(\mathbf{x})$ , i.e.  $I(G) \subseteq I(p)$
- What about the converse? We need one more assumption:
- A distribution is **positive** if  $p(\mathbf{x}) > 0$  for all  $\mathbf{x}$
- Theorem (**Hammersley-Clifford**, 1971): If  $p(\mathbf{x})$  is a positive distribution and  $G$  is an I-map for  $p(\mathbf{x})$ , then  $p(\mathbf{x})$  is a Gibbs distribution that factorizes over  $G$
- Proof is in Koller & Friedman book (as is counter-example for when  $p(\mathbf{x})$  is not positive)
- This is important for **learning**:
  - Prior knowledge is often in the form of conditional independencies (i.e., a graph structure  $G$ )
  - Hammersley-Clifford tells us that it suffices to search over Gibbs distributions for  $G$  – allows us to *parameterize* the distribution



# Today's lecture

- Markov random fields
  - ① Formalism
  - ② Conditional independence properties
  - ③ Examples
- Factor graphs
- Bayesian networks  $\Rightarrow$  Markov random fields (*moralization*)
- Learning graphical models from data

- Let's assume that the domain is governed by some underlying distribution  $p^*$ , which is induced by some network model  $\mathcal{M}^* = (\mathcal{G}^*, \theta^*)$
- We are given a dataset  $\mathcal{D}$  of  $M$  samples from  $p^*$
- The standard assumption is that the data instances are **independent and identically distributed (IID)**
- We are also given a family of models  $\mathcal{M}$ , and our task is to learn some model  $\hat{\mathcal{M}} \in \mathcal{M}$  (i.e., in this family) that defines a distribution  $p_{\hat{\mathcal{M}}}$
- We can learn model parameters for a fixed structure, or both the structure and model parameters

# What should our goal be?

- Return a model  $\hat{\mathcal{M}}$  that precisely captures the distribution  $p^*$  from which our data was sampled
- This is in general not achievable because of
  - computational reasons
  - limited data only provides a rough approximation of the true underlying distribution
- We need to select  $\hat{\mathcal{M}}$  to construct the "best" approximation to  $\mathcal{M}^*$
- What is "best"?

# What is “best”?

This depends on what we want to do

- 1 Density estimation: we are interested in the full distribution (so later we can compute whatever conditional probabilities we want)
- 2 Specific prediction tasks: we are using the distribution to make a prediction  
*We discuss this in depth in lectures 9 and 10.*
- 3 Structure or knowledge discovery: we are interested in the model itself (often of interest in data science)

# Learning for density estimation

- We want to learn the full distribution so that later we can answer *any* probabilistic inference query
- In this setting we can view the learning problem as **density estimation**
- We want to construct  $\hat{\mathcal{M}}$  as “close” as possible to  $p^*$
- How do we evaluate “closeness”?
- **KL-divergence** is one possibility:

$$\mathbf{D}(p^* || p_\theta) = \mathbf{E}_{\mathbf{x} \sim p^*} \left[ \log \left( \frac{p^*(\mathbf{x})}{p_\theta(\mathbf{x})} \right) \right]$$

- $\mathbf{D}(P^* || \hat{P}) = 0$  iff the two distributions are the same.

# Expected log-likelihood

- We can simplify this somewhat:

$$\mathbf{D}(p^* || p_\theta) = \mathbf{E}_{\mathbf{x} \sim p^*} \left[ \log \left( \frac{p^*(\mathbf{x})}{p_\theta(\mathbf{x})} \right) \right] = -\mathbf{H}(p^*) - \mathbf{E}_{\mathbf{x} \sim p^*} [\log p_\theta(\mathbf{x})]$$

- The first term does not depend on  $\theta$ .
- Then, *minimizing* the KL-divergence from the true distribution is equivalent to *maximizing* the **expected log-likelihood**:

$$\mathbf{E}_{\mathbf{x} \sim p^*} [\log p_\theta(\mathbf{x})]$$

- Asks that  $p_\theta$  assign high probability to instances sampled from  $p^*$ , so as to reflect the true distribution
- Because of log, samples  $\mathbf{x}$  where  $p_\theta(\mathbf{x}) \approx 0$  weigh heavily in objective
- **Problem:** In general we do not know  $p^*$ .

# Maximum likelihood

- Approximate the expected log-likelihood

$$\mathbf{E}_{\mathbf{x} \sim p^*} [\log p_{\theta}(\mathbf{x})]$$

with the *empirical log-likelihood*:

$$\mathbf{E}_{\mathcal{D}} [\log p_{\theta}(\mathbf{x})] = \frac{1}{|\mathcal{D}|} \sum_{\mathbf{x} \in \mathcal{D}} \log p_{\theta}(\mathbf{x})$$

- Maximum likelihood learning is then:

$$\max_{\theta} \frac{1}{|\mathcal{D}|} \sum_{\mathbf{x} \in \mathcal{D}} \log p_{\theta}(\mathbf{x})$$

# How to avoid overfitting?

- Hard constraints, e.g. by selecting a less expressive hypothesis class:
  - Bayesian networks with at most  $d$  parents
  - Bayesian networks with parametric conditional probability distributions (e.g. logistic, noisy-or, neural networks)
  - Pairwise MRFs (instead of arbitrary higher-order potentials)
- Augment the learning objective function with **regularization**:

$$\text{objective}(\mathbf{x}, \mathcal{M}) = \text{loss}(\mathbf{x}, \mathcal{M}) + R(\mathcal{M})$$

(often equivalent to MAP estimation where we put a prior over parameters  $\theta$  and maximize  $\log p(\theta \mid \mathbf{x}) = \log p(\mathbf{x}; \theta) + \log p(\theta) - \text{constant}$ )



# ML estimation in Bayesian networks

- Suppose that we know the Bayesian network structure  $G$
- Let  $\theta_{x_i | \mathbf{x}_{pa(i)}}$  be the parameter giving the value of the CPD  $p(x_i | \mathbf{x}_{pa(i)}; \theta)$
- Maximum likelihood estimation corresponds to solving:

$$\max_{\theta} \sum_{n=1}^N \log p(\mathbf{x}^n; \theta) = \max_{\theta} \ell(\theta; \mathcal{D})$$

subject to the non-negativity and normalization constraints

- This is equal to:

$$\begin{aligned} \max_{\theta} \sum_{n=1}^N \log p(\mathbf{x}^n; \theta) &= \max_{\theta} \sum_{n=1}^N \sum_{i=1}^{|V|} \log p(x_i^n | \mathbf{x}_{pa(i)}^n; \theta) \\ &= \max_{\theta} \sum_{i=1}^{|V|} \sum_{n=1}^N \log p(x_i^n | \mathbf{x}_{pa(i)}^n; \theta) \end{aligned}$$

- The optimization problem decomposes into an independent optimization problem for each CPD!

# ML estimation in Bayesian networks

$$\begin{aligned}\ell(\theta; \mathcal{D}) = \log p(\mathcal{D}; \theta) &= \sum_{i=1}^{|V|} \sum_{n=1}^N \log p(x_i^n \mid \mathbf{x}_{pa(i)}^n; \theta) \\ &= \sum_{i=1}^{|V|} \sum_{\mathbf{x}_{pa(i)}} \sum_{\mathbf{x}_i} \sum_{\substack{\hat{\mathbf{x}} \in \mathcal{D}: \\ \hat{x}_i, \hat{\mathbf{x}}_{pa(i)} = \mathbf{x}_i, \mathbf{x}_{pa(i)}}} \log p(x_i \mid \mathbf{x}_{pa(i)}; \theta) \\ &= \sum_{i=1}^{|V|} \sum_{\mathbf{x}_{pa(i)}} \sum_{\mathbf{x}_i} N_{\mathbf{x}_i, \mathbf{x}_{pa(i)}} \log \theta_{\mathbf{x}_i \mid \mathbf{x}_{pa(i)}},\end{aligned}$$

where  $N_{\mathbf{x}_i, \mathbf{x}_{pa(i)}}$  is the number of times that the (partial) assignment  $x_i, \mathbf{x}_{pa(i)}$  is observed in the training data

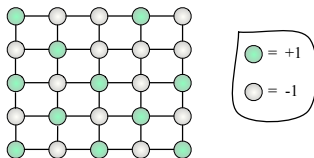
- We have the closed form ML solution:

$$\theta_{\mathbf{x}_i \mid \mathbf{x}_{pa(i)}}^{ML} = \frac{N_{\mathbf{x}_i, \mathbf{x}_{pa(i)}}}{\sum_{\hat{\mathbf{x}}_i} N_{\hat{\mathbf{x}}_i, \mathbf{x}_{pa(i)}}}$$

- We were able to estimate each CPD independently because the objective **decomposes** by variable and parent assignment

# ML estimation in Markov networks

- How do we learn the parameters of an Ising model?



$$p(x_1, \dots, x_n) = \frac{1}{Z} \exp \left( \sum_{i < j} w_{i,j} x_i x_j - \sum_i u_i x_i \right)$$

# Bad news for Markov networks

- The global normalization constant  $Z(\theta)$  kills decomposability:

$$\begin{aligned}\theta^{ML} &= \arg \max_{\theta} \log \prod_{\mathbf{x} \in \mathcal{D}} p(\mathbf{x}; \theta) \\ &= \arg \max_{\theta} \sum_{\mathbf{x} \in \mathcal{D}} \left( \sum_c \log \phi_c(\mathbf{x}_c; \theta) - \log Z(\theta) \right) \\ &= \arg \max_{\theta} \left( \sum_{\mathbf{x} \in \mathcal{D}} \sum_c \log \phi_c(\mathbf{x}_c; \theta) \right) - |\mathcal{D}| \log Z(\theta)\end{aligned}$$

- The log-partition function prevents us from decomposing the objective into a sum over terms for each potential
- Solving for the parameters becomes much more complicated

# Bayesian network structure learning

- Purpose: *knowledge discovery*. We hope that looking at the learned model we can discover something about  $p^*$ , e.g.
  - Nature of the dependencies, e.g., positive or negative correlation
  - What are the direct and indirect dependencies
- Simple statistical models (e.g., looking at correlations) can be used for the first
- But the learned network gives us much more information, e.g. conditional independencies, causal relationships
- In this setting we care about discovering the correct model  $\mathcal{M}^*$ , rather than a different model  $\hat{\mathcal{M}}$  that induces a distribution similar to  $\mathcal{M}^*$ .
- Metric is in terms of the differences between  $\mathcal{M}^*$  and  $\hat{\mathcal{M}}$ .

# This is not always achievable

- The true model might not be identifiable
  - e.g., Bayesian network with several I-equivalent structures.
  - In this case the best we can hope is to discover an I-equivalent structure.
  - Problem is worse when the amount of data is limited and the relationships are weak.

# Structure learning using maximum likelihood

- Recall that for Bayesian networks we have decomposability of the likelihood:

$$\log p(\mathcal{D}; \theta) = \sum_{i=1}^{|V|} \sum_{\mathbf{x}_{pa(i)}} \sum_{\mathbf{x}_i} N_{\mathbf{x}_i, \mathbf{x}_{pa(i)}} \log p(\mathbf{x}_i \mid \mathbf{x}_{pa(i)}; \theta)$$

- Given a candidate structure  $G = (V, E)$ , the maximum likelihood parameters are given by:  $\theta_{\mathbf{x}_i \mid \mathbf{x}_{pa(i)}}^{ML} = \frac{N_{\mathbf{x}_i, \mathbf{x}_{pa(i)}}}{\sum_{\hat{\mathbf{x}}_i} N_{\hat{\mathbf{x}}_i, \mathbf{x}_{pa(i)}}} = \hat{p}(\mathbf{x}_i \mid \mathbf{x}_{pa(i)})$
- Putting these together, maximum likelihood structure learning reduces to:

$$\max_G \sum_{i=1}^{|V|} \text{score}(i \mid pa_i, \mathcal{D}), \quad \text{where}$$

$$\text{score}(i \mid pa_i, \mathcal{D}) = \sum_{\mathbf{x}_{pa(i)}} \sum_{\mathbf{x}_i} N_{\mathbf{x}_i, \mathbf{x}_{pa(i)}} \log p(\mathbf{x}_i \mid \mathbf{x}_{pa(i)}; \theta_{\mathbf{x}_i \mid \mathbf{x}_{pa(i)}}^{ML})$$

# Structure learning using maximum likelihood

- Simplifying further, we get:

$$\begin{aligned}\text{score}(i \mid pa_i, \mathcal{D}) &= N \sum_{\mathbf{x}_{pa(i)}} \frac{N_{\mathbf{x}_{pa(i)}}}{N} \sum_{\mathbf{x}_i} \frac{N_{\mathbf{x}_i, \mathbf{x}_{pa(i)}}}{N_{\mathbf{x}_{pa(i)}}} \log \hat{p}(\mathbf{x}_i \mid \mathbf{x}_{pa(i)}) \\&= N \sum_{\mathbf{x}_{pa(i)}} \hat{p}(\mathbf{x}_{pa(i)}) \sum_{\mathbf{x}_i} \hat{p}(\mathbf{x}_i \mid \mathbf{x}_{pa(i)}) \log \hat{p}(\mathbf{x}_i \mid \mathbf{x}_{pa(i)}) \\&= -N \sum_{\mathbf{x}_{pa(i)}} \hat{p}(\mathbf{x}_{pa(i)}) \sum_{\mathbf{x}_i} \hat{p}(\mathbf{x}_i \mid \mathbf{x}_{pa(i)}) \log \frac{1}{\hat{p}(\mathbf{x}_i \mid \mathbf{x}_{pa(i)})} \\&= -N \cdot \hat{H}(X_i \mid X_{pa(i)}).\end{aligned}$$

- We see that the maximum likelihood structure problem is equivalent to

$$\min_G \sum_{i=1}^N \hat{H}(X_i \mid X_{pa(i)}),$$

i.e. choose a graph structure which minimizes the entropy of each individual variable.



# Structure learning: score-based approaches

- Q: What is the maximum likelihood graph?
  - A: The complete graph! Because  $H(X | Y) \leq H(X)$  **always**.
  - Must *regularize* to recover a sparse graph and have any hope of recovering true structure
  - Common approaches such as BIC and BDe (Bayesian Dirichlet score) are also decomposable
- Obtain a combinatorial optimization problem over acyclic graphs:

$$\text{score}(G; D) = \sum_{i=1}^n \text{score}(i | \text{pa}_i, D)$$


---

**Finding highest scoring graph is NP-hard** – must disallow cycles:



$$\begin{aligned} \text{score}\left(\begin{array}{c} \circ \\ \swarrow \quad \searrow \\ \circ \quad \circ \\ \swarrow \quad \searrow \\ \circ \end{array}\right) &= \text{score}\left(\begin{array}{c} \circ \\ \swarrow \quad \searrow \\ \circ \quad \circ \\ \swarrow \quad \searrow \\ \circ \end{array}\right) + \text{score}(i | \text{pa}_i) \\ &+ \text{score}\left(\begin{array}{c} \circ \\ \swarrow \quad \searrow \\ \circ \quad \circ \\ \swarrow \quad \searrow \\ \circ \end{array}\right) + \text{score}(i | \text{pa}_i) \\ &+ \text{score}\left(\begin{array}{c} \circ \\ \swarrow \quad \searrow \\ \circ \quad \circ \\ \swarrow \quad \searrow \\ \circ \end{array}\right) + \text{score}(i | \text{pa}_i) \\ &+ \text{score}\left(\begin{array}{c} \circ \\ \swarrow \quad \searrow \\ \circ \quad \circ \\ \swarrow \quad \searrow \\ \circ \end{array}\right) + \text{score}(i | \text{pa}_i) \end{aligned}$$

# Independence tests



The network structure implies several conditional independence statements:

$$D \perp I$$

$$G \perp S \mid I$$

$$D \perp L \mid G$$

$$L \perp S \mid G$$

$$L \perp S \mid I$$

$$D \perp S$$

If two variables are (conditionally) independent, structure has no edge between them

- Must make assumption that data is drawn from an I-map of the graph
- Possible to learn structure with polynomial number of data points and polynomial computation time (e.g., the SGS algorithm from Spirtes, Glymour, & Scheines '01)
- Very brittle: if we say that  $X_i \perp X_j \mid X_v$  and they in fact are not, the resulting structure can be very off