

Inference and Representation

David Sontag

New York University

Lecture 3, Sept. 26, 2016

Today's lecture

- Exact inference
 - ① Worst-case complexity of probabilistic inference
 - ② Elimination algorithm
 - ③ Running-time analysis of elimination algorithm (*treewidth*)
- Approximate inference

Probabilistic inference

- Today we consider exact inference in graphical models
- In particular, we focus on conditional probability queries,

$$p(\mathbf{Y}|\mathbf{E} = \mathbf{e}) = \frac{p(\mathbf{Y}, \mathbf{e})}{p(\mathbf{e})}$$

(e.g., the probability of a patient having a disease given some observed symptoms)

- Let $\mathbf{W} = \mathcal{X} - \mathbf{Y} - \mathbf{E}$ be the random variables that are neither the query nor the evidence. Each of these joint distributions can be computed by marginalizing over the other variables:

$$p(\mathbf{Y}, \mathbf{e}) = \sum_{\mathbf{w}} p(\mathbf{Y}, \mathbf{e}, \mathbf{w}), \quad p(\mathbf{e}) = \sum_{\mathbf{y}} p(\mathbf{y}, \mathbf{e})$$

- Naively marginalizing over all unobserved variables requires an exponential number of computations
- Does there exist a more efficient algorithm?

Computational complexity of probabilistic inference

- Here we show that, unless $P=NP$, there does *not* exist a more efficient algorithm
- We show this by reducing 3-SAT, which is NP-hard, to probabilistic inference in Bayesian networks
- 3-SAT asks about the *satisfiability* of a logical formula defined on n literals Q_1, \dots, Q_n , e.g.

$$(\neg Q_3 \vee \neg Q_2 \vee Q_3) \wedge (Q_2 \vee \neg Q_4 \vee \neg Q_5) \dots$$

- Each of the disjunction terms is called a *clause*, e.g.


$$C_1(q_1, q_2, q_3) = \neg q_3 \vee \neg q_2 \vee q_3$$

In 3-SAT, each clause is defined on at most 3 literals.

- Our reduction also proves that inference in Markov networks is NP-hard (why?)

Reducing satisfiability to MAP inference

- **Input:** 3-SAT formula with n literals Q_1, \dots, Q_n and m clauses C_1, \dots, C_m



- One variable $Q_i \in \{0, 1\}$ for each literal, $p(Q_i = 1) = 0.5$.
- One variable $C_i \in \{0, 1\}$ for each clause, whose parents are the literals used in the clause. $C_i = 1$ if the clause is satisfied, and 0 otherwise:


$$p(C_i = 1 \mid \mathbf{q}_{\text{pa}(i)}) = 1[C_i(\mathbf{q}_{\text{pa}(i)})]$$

- Variable X which is 1 if all clauses satisfied, and 0 otherwise:

$$\begin{aligned} p(A_i = 1 \mid \mathbf{pa}(A_i)) &= 1[\mathbf{pa}(A_i) = \mathbf{1}], \text{ for } i = 1, \dots, m-2 \\ p(X = 1 \mid a_{m-2}, c_m) &= 1[a_{m-2} = 1, c_m = 1] \end{aligned}$$

Reducing satisfiability to MAP inference

- **Input:** 3-SAT formula with n literals Q_1, \dots, Q_n and m clauses C_1, \dots, C_m




- $p(\mathbf{q}, \mathbf{c}, \mathbf{a}, X = 1) = 0$ for any assignment \mathbf{q} which does not satisfy all clauses
- $p(\mathbf{Q} = \mathbf{q}, \mathbf{C} = \mathbf{1}, \mathbf{A} = \mathbf{1}, X = 1) = \frac{1}{2^n}$ for any satisfying assignment \mathbf{q}
- Thus, we can find a satisfying assignment (whenever one exists) by constructing this BN and finding the maximum a posteriori (MAP) assignment:

$$\operatorname{argmax}_{\mathbf{q}, \mathbf{c}, \mathbf{a}} p(\mathbf{Q} = \mathbf{q}, \mathbf{C} = \mathbf{c}, \mathbf{A} = \mathbf{a} \mid X = 1)$$

- This proves that MAP inference in Bayesian networks and MRFs is NP-hard

Reducing satisfiability to marginal inference


- **Input:** 3-SAT formula with n literals Q_1, \dots, Q_n and m clauses C_1, \dots, C_m




- $p(X = 1) = \sum_{\mathbf{q}, \mathbf{c}, \mathbf{a}} p(\mathbf{Q} = \mathbf{q}, \mathbf{C} = \mathbf{c}, \mathbf{A} = \mathbf{a}, X = 1)$ is equal to the number of satisfying assignments times $\frac{1}{2^n}$
- Thus, $p(X = 1) > 0$ if and only if the formula has a satisfying assignment
- This shows that *marginal inference* is also NP-hard

Probabilistic inference in practice

- NP-hardness simply says that there **exist** difficult inference problems
- Real-world inference problems are not necessarily as hard as these worst-case instances
- The reduction from SAT created a very complex Bayesian network:



- Some graphs are **easy** to do inference in! For example, inference in hidden Markov models




and other tree-structured graphs can be performed in **linear time**

Marginal inference in HMMs

- “Filtering” problem is to do marginal inference to find:

$$\Pr(x_n \mid y_1, \dots, y_n)$$



- How does one **compute** this?
- Applying rule of conditional probability, we have:


$$\Pr(x_n \mid y_1, \dots, y_n) = \frac{\Pr(x_n, y_1, \dots, y_n)}{\Pr(y_1, \dots, y_n)}$$

- Naively, would seem to require k^{n-1} summations,

$$\Pr(x_n, y_1, \dots, y_n) = \sum_{x_1, \dots, x_{n-1}} \Pr(x_1, \dots, x_n, y_1, \dots, y_n)$$

Is there a
more efficient
algorithm?

Marginal inference in HMMs:



- Use **dynamic programming**

$$\begin{aligned}
 \Pr(x_n, y_1, \dots, y_n) &= \sum_{x_{n-1}} \Pr(x_{n-1}, x_n, y_1, \dots, y_n) \\
 &= \sum_{x_{n-1}} \Pr(x_{n-1}, y_1, \dots, y_{n-1}) \Pr(x_n, y_n \mid x_{n-1}, y_1, \dots, y_{n-1}) \\
 &= \sum_{x_{n-1}} \Pr(x_{n-1}, y_1, \dots, y_{n-1}) \Pr(x_n, y_n \mid x_{n-1}) \\
 &= \sum_{x_{n-1}} \Pr(x_{n-1}, y_1, \dots, y_{n-1}) \Pr(x_n \mid x_{n-1}) \Pr(y_n \mid x_n, x_{n-1}) \\
 &= \sum_{x_{n-1}} \Pr(x_{n-1}, y_1, \dots, y_{n-1}) \Pr(x_n \mid x_{n-1}) \Pr(y_n \mid x_n)
 \end{aligned}$$

$$\Pr(A = a) = \sum_b \Pr(B = b, A = a)$$

$$\Pr(\vec{A} = \vec{a}, \vec{B} = \vec{b}) = \Pr(\vec{A} = \vec{a}) \Pr(\vec{B} = \vec{b} \mid \vec{A} = \vec{a})$$


Conditional independence in HMMs

$$\Pr(A = a, B = b) = \Pr(A = a) \Pr(B = b \mid A = a)$$

Conditional independence in HMMs

- For $n=1$, initialize $\Pr(x_1, y_1) = \Pr(x_1) \Pr(y_1 \mid x_1)$
- Total running time is $O(nk^2)$ – linear time! **Easy to do filtering**


MAP inference in HMMs:



- MAP inference in HMMs can be solved in linear time!

$$\begin{aligned} \arg \max_{\mathbf{x}} \Pr(x_1, \dots, x_n \mid y_1, \dots, y_n) &= \arg \max_{\mathbf{x}} \Pr(x_1, \dots, x_n, y_1, \dots, y_n) \\ &= \arg \max_{\mathbf{x}} \log \Pr(x_1, \dots, x_n, y_1, \dots, y_n) \\ &= \arg \max_{\mathbf{x}} \log \left[\Pr(x_1) \Pr(y_1 \mid x_1) \right] + \sum_{i=2}^n \log \left[\Pr(x_i \mid x_{i-1}) \Pr(y_i \mid x_i) \right] \end{aligned}$$

- Formulate as a shortest paths problem



Called the Viterbi algorithm

Variable elimination (VE)

- Exact algorithm for probabilistic inference in **any** graphical model
- Running time will depend on the *graph structure*
- Uses **dynamic programming** to circumvent enumerating all assignments
- First we introduce the concept for computing marginal probabilities, $p(X_i)$, in Bayesian networks
- After this, we will generalize to MRFs and conditional queries

- Suppose we have a simple chain, $A \rightarrow B \rightarrow C \rightarrow D$, and we want to compute $p(D)$
- $p(D)$ is a **set** of values, $\{p(D = d), d \in \text{Val}(D)\}$. Algorithm computes sets of values at a time – an entire distribution
- By the chain rule and conditional independence, the joint distribution factors as

$$p(A, B, C, D) = p(A)p(B | A)p(C | B)p(D | C)$$

- In order to compute $p(D)$, we have to marginalize over A, B, C :

$$p(D) = \sum_{a,b,c} p(A = a, B = b, C = c, D)$$

Let's be a bit more explicit...

$$\begin{array}{cccc}
 & P(a^1) & P(b^1 | a^1) & P(c^1 | b^1) & P(d^1 | c^1) \\
 + & P(a^2) & P(b^1 | a^2) & P(c^1 | b^1) & P(d^1 | c^1) \\
 + & P(a^1) & P(b^2 | a^1) & P(c^1 | b^2) & P(d^1 | c^1) \\
 + & P(a^2) & P(b^2 | a^2) & P(c^1 | b^2) & P(d^1 | c^1) \\
 + & P(a^1) & P(b^1 | a^1) & P(c^2 | b^1) & P(d^1 | c^2) \\
 + & P(a^2) & P(b^1 | a^2) & P(c^2 | b^1) & P(d^1 | c^2) \\
 + & P(a^1) & P(b^2 | a^1) & P(c^2 | b^2) & P(d^1 | c^2) \\
 + & P(a^2) & P(b^2 | a^2) & P(c^2 | b^2) & P(d^1 | c^2)
 \end{array}$$

$$\begin{array}{cccc}
 & P(a^1) & P(b^1 | a^1) & P(c^1 | b^1) & P(d^2 | c^1) \\
 + & P(a^2) & P(b^1 | a^2) & P(c^1 | b^1) & P(d^2 | c^1) \\
 + & P(a^1) & P(b^2 | a^1) & P(c^1 | b^2) & P(d^2 | c^1) \\
 + & P(a^2) & P(b^2 | a^2) & P(c^1 | b^2) & P(d^2 | c^1) \\
 + & P(a^1) & P(b^1 | a^1) & P(c^2 | b^1) & P(d^2 | c^2) \\
 + & P(a^2) & P(b^1 | a^2) & P(c^2 | b^1) & P(d^2 | c^2) \\
 + & P(a^1) & P(b^2 | a^1) & P(c^2 | b^2) & P(d^2 | c^2) \\
 + & P(a^2) & P(b^2 | a^2) & P(c^2 | b^2) & P(d^2 | c^2)
 \end{array}$$

- There is structure to the summation, e.g., repeated $P(c^1|b^1)P(d^1|c^1)$

Simplifying the computation

- Our goal was to compute

$$\begin{aligned} p(D) &= \sum_{a,b,c} p(a, b, c, D) = \sum_{a,b,c} p(a)p(b|a)p(c|b)p(D|c) \\ &= \sum_c \sum_b \sum_a p(D|c)p(c|b)p(b|a)p(a) \end{aligned}$$

- We can push the summations inside to obtain:

$$p(D) = \sum_c p(D|c) \sum_b p(c|b) \underbrace{\sum_a \underbrace{p(b|a)p(a)}_{\psi_1(a,b)}}_{\tau_1(b)}$$

- Let's call $\psi_1(A, B) = P(A)P(B|A)$. Then, $\tau_1(B) = \sum_a \psi_1(a, B)$
- Similarly, let $\psi_2(B, C) = \tau_1(B)P(C|B)$. Then, $\tau_2(C) = \sum_b \psi_2(b, C)$
- This procedure is dynamic programming: computation is inside out instead of outside in

Summary so far

- Worst-case analysis says that marginal inference is NP-hard
- In practice, due to the structure of the Bayesian network, we can cache computations that are otherwise computed exponentially many times
- This depends on our having a good **variable elimination ordering**

Sum-product inference task

- We want to give an algorithm to compute $p(\mathbf{Y})$ for BNs and MRFs
- This can be reduced to the following **sum-product** inference task:

$$\text{Compute } \tau(\mathbf{y}) = \sum_{\mathbf{z}} \prod_{\phi \in \Phi} \phi(\mathbf{z}_{\text{Scope}[\phi] \cap \mathbf{Z}}, \mathbf{y}_{\text{Scope}[\phi] \cap \mathbf{Y}}) \quad \forall \mathbf{y},$$

where Φ is a set of factors or potentials

- For a BN, Φ is given by the conditional probability distributions for all variables,

$$\Phi = \{\phi_{X_i}\}_{i=1}^n = \{p(X_i \mid \mathbf{X}_{\text{Pa}(X_i)})\}_{i=1}^n,$$

and where we sum over the set $\mathbf{Z} = \mathcal{X} - \mathbf{Y}$

- For Markov networks, the factors Φ correspond to the set of potentials which we earlier called \mathcal{C}
 - Sum-product returns an unnormalized distribution, so we divide by $\sum_{\mathbf{y}} \tau(\mathbf{y})$

Factor marginalization

- Let $\phi(\mathbf{X}, Y)$ be a factor where \mathbf{X} is a set of variables and $Y \notin \mathbf{X}$
- Factor marginalization** of ϕ over Y (also called “summing out Y in ϕ ”) gives a new factor:

$$\tau(\mathbf{X}) = \sum_Y \phi(\mathbf{X}, Y)$$

For example,

a^1	b^1	c^1	0.25
a^1	b^1	c^2	0.35
a^1	b^2	c^1	0.08
a^1	b^2	c^2	0.16
a^2	b^1	c^1	0.05
a^2	b^1	c^2	0.07
a^2	b^2	c^1	0
a^2	b^2	c^2	0
a^3	b^1	c^1	0.15
a^3	b^1	c^2	0.21
a^3	b^2	c^1	0.09
a^3	b^2	c^2	0.18

a^1	c^1	0.33
a^1	c^2	0.51
a^2	c^1	0.05
a^2	c^2	0.07
a^3	c^1	0.24
a^3	c^2	0.39

Sum-product variable elimination

- Order the variables \mathbf{Z} (called the **elimination ordering**)
- Iteratively marginalize out variable Z_i , one at a time
- For each i ,
 - 1 Multiply all factors that have Z_i in their scope, generating a new product factor
 - 2 Marginalize this product factor over Z_i , generating a smaller factor
 - 3 Remove the old factors from the set of all factors, and add the new one

Algorithm 9.1 Sum-Product Variable Elimination algorithm


Procedure Sum-Product-Variable-Elimination (
 Φ , // Set of factors
 Z , // Set of variables to be eliminated
 \prec // Ordering on Z
)

1 Let Z_1, \dots, Z_k be an ordering of Z such that
2 $Z_i \prec Z_j$ iff $i < j$
3 **for** $i = 1, \dots, k$
4 $\Phi \leftarrow$ Sum-Product-Eliminate-Var(Φ, Z_i)
5 $\phi^* \leftarrow \prod_{\phi \in \Phi} \phi$
6 **return** ϕ^*

Procedure Sum-Product-Eliminate-Var (
 Φ , // Set of factors
 Z // Variable to be eliminated
)

1 $\Phi' \leftarrow \{\phi \in \Phi : Z \in \text{Scope}[\phi]\}$
2 $\Phi'' \leftarrow \Phi - \Phi'$
3 $\psi \leftarrow \prod_{\phi \in \Phi'} \phi$
4 $\tau \leftarrow \sum_Z \psi$
5 **return** $\Phi'' \cup \{\tau\}$

Example



- What is $p(\text{Job})$? Joint distribution factorizes as:

$$p(C, D, I, G, S, L, H, J) = p(C)p(D|C)p(I)p(G|D, I)p(L|G)P(S|I)P(J|S, L)p(H|J, G)$$


with factors

$$\Phi = \{\phi_C(C), \phi_D(C, D), \phi_I(I), \phi_G(G, D, I), \phi_L(L, G), \\ \phi_S(S, I), \phi_J(J, S, L), \phi_H(H, J, G)\}$$

- Let's do variable elimination with ordering $\{C, D, I, H, G, S, L\}$ on the board!

Elimination ordering

- We can pick any order we want, but some orderings introduce factors with much larger scope



Step	Variable eliminated	Factors used	Variables involved	New factor
1	C	$\phi_C(C), \phi_D(D, C)$	C, D	$\tau_1(D)$
2	D	$\phi_G(G, I, D), \tau_1(D)$	G, I, D	$\tau_2(G, I)$
3	I	$\phi_I(I), \phi_S(S, I), \tau_2(G, I)$	G, S, I	$\tau_3(G, S)$
4	H	$\phi_H(H, G, J)$	H, G, J	$\tau_4(G, J)$
5	G	$\tau_4(G, J), \tau_3(G, S), \phi_L(L, G)$	G, J, L, S	$\tau_5(J, L, S)$
6	S	$\tau_5(J, L, S), \phi_J(J, L, S)$	J, L, S	$\tau_6(J, L)$
7	L	$\tau_6(J, L)$	J, L	$\tau_7(J)$

- Alternative ordering...

Step	Variable eliminated	Factors used	Variables involved	New factor
1	G	$\phi_G(G, I, D), \phi_L(L, G), \phi_H(H, G, J)$	G, I, D, L, J, H	$\tau_1(I, D, L, J, H)$
2	I	$\phi_I(I), \phi_S(S, I), \tau_1(I, D, L, S, J, H)$	S, I, D, L, J, H	$\tau_2(D, L, S, J, H)$
3	S	$\phi_J(J, L, S), \tau_2(D, L, S, J, H)$	D, L, S, J, H	$\tau_3(D, L, J, H)$
4	L	$\tau_3(D, L, J, H)$	D, L, J, H	$\tau_4(D, J, H)$
5	H	$\tau_4(D, J, H)$	D, J, H	$\tau_5(D, J)$
6	C	$\tau_5(D, J), \phi_D(D, C)$	D, J, C	$\tau_6(D, J)$
7	D	$\tau_6(D, J)$	D, J	$\tau_7(J)$

How to introduce evidence?

- Recall that our original goal was to answer *conditional* probability queries,


$$p(\mathbf{Y} | \mathbf{E} = \mathbf{e}) = \frac{p(\mathbf{Y}, \mathbf{e})}{p(\mathbf{e})}$$

- Apply variable elimination algorithm to the task of computing $P(\mathbf{Y}, \mathbf{e})$
- Replace each factor $\phi \in \Phi$ that has $\mathbf{E} \cap \text{Scope}[\phi] \neq \emptyset$ with

$$\phi'(\mathbf{x}_{\text{Scope}[\phi] - \mathbf{E}}) = \phi(\mathbf{x}_{\text{Scope}[\phi] - \mathbf{E}}, \mathbf{e}_{\mathbf{E} \cap \text{Scope}[\phi]})$$

- Then, eliminate the variables in $\mathcal{X} - \mathbf{Y} - \mathbf{E}$. The returned factor $\phi^*(\mathbf{Y})$ is $p(\mathbf{Y}, \mathbf{e})$
- To obtain the conditional $p(\mathbf{Y} | \mathbf{e})$, normalize the resulting product of factors – the normalization constant is $p(\mathbf{e})$

Running time of variable elimination





Step	Variable eliminated	Factors used	Variables involved	New factor
1	C	$\phi_C(C), \phi_D(D, C)$	C, D	$\tau_1(D)$
2	D	$\phi_G(G, I, D), \tau_1(D)$	G, I, D	$\tau_2(G, I)$
3	I	$\phi_I(I), \phi_S(S, I), \tau_2(G, I)$	G, S, I	$\tau_3(G, S)$
4	H	$\phi_H(H, G, J)$	H, G, J	$\tau_4(G, J)$
5	G	$\tau_4(G, J), \tau_3(G, S), \phi_L(L, G)$	G, J, L, S	$\tau_5(J, L, S)$
6	S	$\tau_5(J, L, S), \phi_J(J, L, S)$	J, L, S	$\tau_6(J, L)$
7	L	$\tau_6(J, L)$	J, L	$\tau_7(J)$

- Let n be the number of variables, and m the number of initial factors
- At each step, we pick a variable X_i and multiply all factors involving X_i , resulting in a single factor ψ_i
- Let N_i be the number of variables in the factor ψ_i , and let $N_{\max} = \max_i N_i$
- The running time of VE is then $O(mk^{N_{\max}})$, where $k = |\text{Val}(X)|$. Why?
- The primary concern is that N_{\max} can potentially be as large as n

Running time in graph-theoretic concepts

- Let's try to analyze the complexity in terms of the graph structure
- G_ϕ is the undirected graph with one node per variable, where there is an edge (X_i, X_j) if these appear together in the scope of some factor ϕ
- Ignoring evidence, this is either the original MRF (for sum-product VE on MRFs) or the moralized Bayesian network:



Elimination as graph transformation


When a variable X is eliminated,

- We create a single factor ψ that contains X and all of the variables \mathbf{Y} with which it appears in factors
- We eliminate X from ψ , replacing it with a new factor τ that contains all of the variables \mathbf{Y} , but not X . Let's call the new set of factors Φ_X


How does this modify the graph, going from G_Φ to G_{Φ_X} ?

- Constructing ψ generates edges between all of the variables $Y \in \mathbf{Y}$
- Some of these edges were already in G_Φ , some are new
- The new edges are called **fill edges**
- The step of removing X from Φ to construct Φ_X removes X and all its incident edges from the graph


Example




(Graph)



(Elim. C)




(Elim. D)




(Elim. I)

- We can summarize the computation cost using a single graph that is the union of all the graphs resulting from each step of the elimination
- We call this the **induced graph** $\mathcal{I}_{\Phi, \prec}$, where \prec is the elimination ordering


Example



Step	Variable eliminated	Factors used	Variables involved	New factor
1	C	$\phi_C(C), \phi_D(D, C)$	C, D	$\tau_1(D)$
2	D	$\phi_G(G, I, D), \tau_1(D)$	G, I, D	$\tau_2(G, I)$
3	I	$\phi_I(I), \phi_S(S, I), \tau_2(G, I)$	G, S, I	$\tau_3(G, S)$
4	H	$\phi_H(H, G, J)$	H, G, J	$\tau_4(G, J)$
5	G	$\tau_4(G, J), \tau_3(G, S), \phi_L(L, G)$	G, J, L, S	$\tau_5(J, L, S)$
6	S	$\tau_5(J, L, S), \phi_J(J, L, S)$	J, L, S	$\tau_6(J, L)$
7	L	$\tau_6(J, L)$	J, L	$\tau_7(J)$



(Induced graph)



(Maximal Cliques)

Properties of the induced graph

- **Theorem:** Let $\mathcal{I}_{\Phi, \prec}$ be the induced graph for a set of factors Φ and ordering \prec , then
 - 1 Every factor generated during VE has a scope that is a clique in $\mathcal{I}_{\Phi, \prec}$
 - 2 Every maximal clique in $\mathcal{I}_{\Phi, \prec}$ is the scope of some intermediate factor in the computation(see Koller & Friedman for proof)
- Thus, N_{max} is equal to the size of the largest clique in $\mathcal{I}_{\Phi, \prec}$
- The running time, $O(mk^{N_{max}})$, is exponential in the size of the largest clique of the induced graph

Induced width

- The **width** of an induced graph is #nodes in largest clique - 1
- We define the **induced width** $w_{\mathcal{G}, \prec}$ to be the width of the graph $\mathcal{I}_{\mathcal{G}, \prec}$ induced by applying VE to \mathcal{G} using ordering \prec
- The **treewidth**, or “minimal induced width” of graph \mathcal{G} is

$$w_{\mathcal{G}}^* = \min_{\prec} w_{\mathcal{G}, \prec}$$

- The treewidth provides a bound on the best running time achievable by VE on a distribution that factorizes over \mathcal{G} : $O(mk^{w_{\mathcal{G}}^*+1})$,
- Unfortunately, finding the **best** elimination ordering (equivalently, computing the treewidth) for a graph is NP-hard
- In practice, heuristics are used to find a good elimination ordering

Today's lecture

- Exact inference
 - ① Worst-case complexity of probabilistic inference
 - ② Elimination algorithm
 - ③ Running-time analysis of elimination algorithm (*treewidth*)
- Approximate inference

Approximate marginal inference

- Given the joint $p(x_1, \dots, x_n)$ represented as a graphical model, how do we perform **marginal inference**, e.g. to compute $p(x_1 \mid e)$?
- We showed earlier in the lecture that doing this exactly is NP-hard
- Nearly all *approximate inference* algorithms are either:
 - 1 Monte-carlo methods (e.g., **Gibbs sampling**, likelihood reweighting, MCMC)
 - 2 Variational algorithms (e.g., mean-field, loopy belief propagation)

Generating samples from a Bayesian network

Algorithm 12.1 Forward Sampling in a Bayesian network

Procedure Forward-Sample (
 \mathcal{B} // Bayesian network over \mathcal{X}
)

1 Let X_1, \dots, X_n be a topological ordering of \mathcal{X}
2 **for** $i = 1, \dots, n$
3 $\mathbf{u}_i \leftarrow \mathbf{x} \langle \text{Pa}_{X_i} \rangle$ // Assignment to Pa_{X_i} in x_1, \dots, x_{i-1}
4 Sample x_i from $P(X_i \mid \mathbf{u}_i)$
5 **return** (x_1, \dots, x_n)

(Koller & Friedman, *Probabilistic Graphical Models*, MIT Press 2009)

Monte-Carlo algorithms

- Given a joint distribution $p(x_1, \dots, x_n)$, how do we compute marginals?

$$\begin{aligned} p[X_1 = x_1] &= E_{\mathbf{x} \sim p}[f(\mathbf{x})], \text{ where } f(\mathbf{x}) = 1[X_1 = x_1] \\ &= \sum_{\mathbf{x}} p(\mathbf{x}) f(\mathbf{x}). \end{aligned}$$

- Rather than explicitly enumerating *all* assignments, consider the following Monte-Carlo estimate of the expectation:

$$\begin{aligned} \mathbf{x}^1 &\sim p(\mathbf{x}) \\ \mathbf{x}^2 &\sim p(\mathbf{x}) \\ &\vdots \\ \mathbf{x}^M &\sim p(\mathbf{x}) \end{aligned}$$

- Then, our *estimate* is $\hat{E}_p[f(\mathbf{x})] = \frac{1}{M} \sum_{m=1}^M f(\mathbf{x}^m)$. **How good is it?**