

# Inference and Representation

David Sontag

New York University

Lecture 10, Nov. 28, 2016

# Reminder: conditional random fields (CRFs)

- A CRF is a Markov network on variables  $\mathbf{X} \cup \mathbf{Y}$ , which specifies the conditional distribution

$$P(\mathbf{y} \mid \mathbf{x}) = \frac{1}{Z(\mathbf{x})} \prod_{c \in C} \phi_c(\mathbf{y}_c, \mathbf{x})$$

with partition function

$$Z(\mathbf{x}) = \sum_{\hat{\mathbf{y}}} \prod_{c \in C} \phi_c(\hat{\mathbf{y}}_c, \mathbf{x}).$$

- As before, two variables in the graph are connected with an undirected edge if they appear together in the scope of some factor
- The only difference with a standard Markov network is the normalization term – before marginalized over  $\mathbf{X}$  and  $\mathbf{Y}$ , now only over  $\mathbf{Y}$

# Reminder: log-linear parameterization of CRFs

- Parameterize  $\phi_c(\mathbf{y}_c, \mathbf{x})$  using a log-linear parameterization:
  - Single weight vector  $\mathbf{w} \in \mathbb{R}^d$  that is used globally
  - For each potential  $c$ , a vector-valued **feature function**  $\mathbf{f}_c(\mathbf{y}_c, \mathbf{x}) \in \mathbb{R}^d$
  - Then,  $\phi_c(\mathbf{y}_c, \mathbf{x}; \mathbf{w}) = \exp(\mathbf{w} \cdot \mathbf{f}_c(\mathbf{y}_c, \mathbf{x}))$
- The conditional distribution is in the *exponential family*!

$$p(\mathbf{y} \mid \mathbf{x}; \mathbf{w}) = \exp\{\mathbf{w} \cdot \mathbf{f}(\mathbf{y}, \mathbf{x}) - \log Z(\mathbf{x}; \mathbf{w})\},$$

where  $\mathbf{f}(\mathbf{y}, \mathbf{x}) = \sum_c \mathbf{f}_c(\mathbf{y}_c, \mathbf{x})$  and  $Z(\mathbf{x}; \mathbf{w}) = \sum_{\hat{\mathbf{y}}} \exp\{\mathbf{w} \cdot \mathbf{f}(\hat{\mathbf{y}}, \mathbf{x})\}$

- This formulation allows for parameter sharing

## Reminder: density estimation for CRFs

$$\text{CRF: } p(\mathbf{y} \mid \mathbf{x}) = \frac{1}{Z(\mathbf{x})} \prod_{c \in C} \phi_c(\mathbf{y}_c, \mathbf{x}), \quad Z(\mathbf{x}) = \sum_{\hat{\mathbf{y}}} \prod_{c \in C} \phi_c(\hat{\mathbf{y}}_c, \mathbf{x})$$

- Empirical risk minimization with CRFs, i.e.  $\min_{\hat{\mathcal{M}}} \mathbf{E}_{\mathcal{D}} [\text{loss}(\mathbf{x}, \mathbf{y}, \hat{\mathcal{M}})]$ :

$$\begin{aligned} \mathbf{w}^{ML} &= \arg \min_{\mathbf{w}} \frac{1}{|\mathcal{D}|} \sum_{(\mathbf{x}, \mathbf{y}) \in \mathcal{D}} -\log p(\mathbf{y} \mid \mathbf{x}; \mathbf{w}) \\ &= \arg \max_{\mathbf{w}} \sum_{(\mathbf{x}, \mathbf{y}) \in \mathcal{D}} \left( \sum_c \log \phi_c(\mathbf{y}_c, \mathbf{x}; \mathbf{w}) - \log Z(\mathbf{x}; \mathbf{w}) \right) \\ &= \arg \max_{\mathbf{w}} \mathbf{w} \cdot \left( \sum_{(\mathbf{x}, \mathbf{y}) \in \mathcal{D}} \sum_c \mathbf{f}_c(\mathbf{y}_c, \mathbf{x}) \right) - \sum_{(\mathbf{x}, \mathbf{y}) \in \mathcal{D}} \log Z(\mathbf{x}; \mathbf{w}) \end{aligned}$$

- What if prediction is only done with MAP inference? Then, the partition function is irrelevant. Is there a way to train to take advantage of this?

# Goal of learning

- The goal of learning is to return a model  $\hat{\mathcal{M}}$  that precisely captures the distribution  $p^*$  from which our data was sampled
- This is in general not achievable because of
  - computational reasons
  - limited data only provides a rough approximation of the true underlying distribution
- We need to select  $\hat{\mathcal{M}}$  to construct the "best" approximation to  $\mathcal{M}^*$
- What is "best"?

# What notion of “best” should learning be optimizing?

This depends on what we want to do

- ① Density estimation: we are interested in the full distribution (so later we can compute whatever conditional probabilities we want)
- ② Specific prediction tasks: we are using the distribution to make a prediction
- ③ Structure or knowledge discovery: we are interested in the model itself

# Structured prediction

- Often we learn a model for the purpose of **structured prediction**, in which given  $\mathbf{x}$  we predict  $\mathbf{y}$  by finding the MAP assignment:

$$\operatorname{argmax}_{\mathbf{y}} \hat{p}(\mathbf{y}|\mathbf{x})$$

- Rather than learn using log-loss (density estimation), we use a loss function better suited to the specific task
- One reasonable choice would be the **classification error**:

$$\mathbf{E}_{(\mathbf{x}, \mathbf{y}) \sim p^*} [\mathbb{1}\{ \exists \mathbf{y}' \neq \mathbf{y} \text{ s.t. } \hat{p}(\mathbf{y}'|\mathbf{x}) \geq \hat{p}(\mathbf{y}|\mathbf{x}) \}]$$

which is the probability over all  $(\mathbf{x}, \mathbf{y})$  pairs sampled from  $p^*$  that our classifier selects the right labels

- If  $p^*$  is in the model family, training with log-loss (density estimation) and classification error would perform similarly (given sufficient data)
- Otherwise, better to directly go for what we care about (classification error)

# Structured prediction

- Consider the empirical risk for 0-1 loss (classification error):

$$\frac{1}{|\mathcal{D}|} \sum_{(\mathbf{x}, \mathbf{y}) \in \mathcal{D}} \mathbb{1}\{ \exists \mathbf{y}' \neq \mathbf{y} \text{ s.t. } \hat{p}(\mathbf{y}'|\mathbf{x}) \geq \hat{p}(\mathbf{y}|\mathbf{x}) \}$$

- Each constraint  $\hat{p}(\mathbf{y}'|\mathbf{x}) \geq \hat{p}(\mathbf{y}|\mathbf{x})$  is equivalent to

$$\mathbf{w} \cdot \sum_c \mathbf{f}_c(\mathbf{x}, \mathbf{y}'_c) - \log Z(\mathbf{x}; \mathbf{w}) \geq \mathbf{w} \cdot \sum_c \mathbf{f}_c(\mathbf{x}, \mathbf{y}_c) - \log Z(\mathbf{x}; \mathbf{w})$$

- The log-partition function cancels out on both sides. Re-arranging, we have:

$$\mathbf{w} \cdot \left( \sum_c \mathbf{f}_c(\mathbf{x}, \mathbf{y}'_c) - \sum_c \mathbf{f}_c(\mathbf{x}, \mathbf{y}_c) \right) \geq 0$$

- Said differently, the empirical risk is **zero** when  $\forall (\mathbf{x}, \mathbf{y}) \in \mathcal{D}$  and  $\mathbf{y}' \neq \mathbf{y}$ ,

$$\mathbf{w} \cdot \left( \sum_c \mathbf{f}_c(\mathbf{x}, \mathbf{y}_c) - \sum_c \mathbf{f}_c(\mathbf{x}, \mathbf{y}'_c) \right) > 0.$$



# Structured prediction

- Empirical risk is **zero** when  $\forall (\mathbf{x}, \mathbf{y}) \in \mathcal{D}$  and  $\mathbf{y}' \neq \mathbf{y}$ ,

$$\mathbf{w} \cdot \left( \sum_c \mathbf{f}_c(\mathbf{x}, \mathbf{y}_c) - \sum_c \mathbf{f}_c(\mathbf{x}, \mathbf{y}'_c) \right) > 0.$$

- In the simplest setting, learning corresponds to finding a weight vector  $\mathbf{w}$  that satisfies all of these constraints (when possible)
- This is a linear program (LP)!
- How many constraints does it have?  $|\mathcal{D}| * |\mathcal{Y}|$  – exponentially many!
- Thus, we must avoid explicitly representing this LP
- The first part of this lecture is about algorithms for solving this LP (or some variant) in a tractable manner

# Structured *perceptron* algorithm

- **Input:** Training examples  $\mathcal{D} = \{(\mathbf{x}^m, \mathbf{y}^m)\}$
- Let  $\mathbf{f}(\mathbf{x}, \mathbf{y}) = \sum_c \mathbf{f}_c(\mathbf{x}, \mathbf{y}_c)$ . Then, the constraints that we want to satisfy are

$$\mathbf{w} \cdot \left( \mathbf{f}(\mathbf{x}^m, \mathbf{y}^m) - \mathbf{f}(\mathbf{x}^m, \mathbf{y}) \right) > 0, \quad \forall \mathbf{y} \neq \mathbf{y}^m$$

- The perceptron algorithm uses MAP inference in its inner loop:

$$\text{MAP}(\mathbf{x}^m; \mathbf{w}) = \arg \max_{\mathbf{y} \in \mathcal{Y}} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}^m, \mathbf{y})$$

The maximization can often be performed efficiently by using the structure!

- The perceptron algorithm is then:
  - 1 Start with  $\mathbf{w} = 0$
  - 2 While the weight vector is still changing:
    - 3 For  $m = 1, \dots, |\mathcal{D}|$
    - 4  $\mathbf{y} \leftarrow \text{MAP}(\mathbf{x}^m; \mathbf{w})$
    - 5  $\mathbf{w} \leftarrow \mathbf{w} + \mathbf{f}(\mathbf{x}^m, \mathbf{y}^m) - \mathbf{f}(\mathbf{x}^m, \mathbf{y})$

# Structured perceptron algorithm

- If the training data is *separable*, the perceptron algorithm is guaranteed to find a weight vector which perfectly classifies all of the data
- When separable with margin  $\gamma$ , number of iterations is at most

$$\left(\frac{2R}{\gamma}\right)^2,$$

where  $R = \max_{m,y} \|\mathbf{f}(\mathbf{x}^m, \mathbf{y})\|_2$

- In practice, one stops after a certain number of outer iterations (called *epochs*), and uses the *average* of all weights
- The averaging can be understood as a type of regularization to prevent overfitting

# Allowing slack

- We can equivalently write the constraints as

$$\mathbf{w} \cdot \left( \mathbf{f}(\mathbf{x}^m, \mathbf{y}^m) - \mathbf{f}(\mathbf{x}^m, \mathbf{y}) \right) \geq 1, \quad \forall \mathbf{y} \neq \mathbf{y}^m$$

- Suppose there do not exist weights  $\mathbf{w}$  that satisfy all constraints
- Introduce *slack* variables  $\xi_m \geq 0$ , one per data point, to allow for constraint violations:

$$\mathbf{w} \cdot \left( \mathbf{f}(\mathbf{x}^m, \mathbf{y}^m) - \mathbf{f}(\mathbf{x}^m, \mathbf{y}) \right) \geq 1 - \xi_m, \quad \forall \mathbf{y} \neq \mathbf{y}^m$$

- Then, minimize the sum of the slack variables,  $\min_{\xi \geq 0} \sum_m \xi_m$ , subject to the above constraints

# Structural SVM (support vector machine)

$$\min_{\mathbf{w}, \xi} \sum_m \xi_m + C \|\mathbf{w}\|^2$$

subject to:

$$\begin{aligned} \mathbf{w} \cdot \left( \mathbf{f}(\mathbf{x}^m, \mathbf{y}^m) - \mathbf{f}(\mathbf{x}^m, \mathbf{y}) \right) &\geq 1 - \xi_m, \quad \forall m, \mathbf{y} \neq \mathbf{y}^m \\ \xi_m &\geq 0, \quad \forall m \end{aligned}$$

This is a quadratic program (QP). Solving for the slack variables in closed form, we obtain

$$\xi_m^* = \max \left( 0, \max_{\mathbf{y} \in \mathcal{Y} \setminus \mathbf{y}^m} 1 - \mathbf{w} \cdot \left( \mathbf{f}(\mathbf{x}^m, \mathbf{y}^m) - \mathbf{f}(\mathbf{x}^m, \mathbf{y}) \right) \right)$$

Thus, we can re-write the whole optimization problem as

$$\min_{\mathbf{w}} \sum_m \max \left( 0, \max_{\mathbf{y} \in \mathcal{Y} \setminus \mathbf{y}^m} 1 - \mathbf{w} \cdot \left( \mathbf{f}(\mathbf{x}^m, \mathbf{y}^m) - \mathbf{f}(\mathbf{x}^m, \mathbf{y}) \right) \right) + C \|\mathbf{w}\|^2$$

# Hinge loss

- We can view  $\max\left(0, \max_{\mathbf{y} \in \mathcal{Y} \setminus \mathbf{y}^m} 1 - \mathbf{w} \cdot \left(\mathbf{f}(\mathbf{x}^m, \mathbf{y}^m) - \mathbf{f}(\mathbf{x}^m, \mathbf{y})\right)\right)$  as a loss function, called *hinge loss*
- When  $\mathbf{w} \cdot \mathbf{f}(\mathbf{x}^m, \mathbf{y}^m) \geq \mathbf{w} \cdot \mathbf{f}(\mathbf{x}^m, \mathbf{y})$  for all  $\mathbf{y}$  (i.e., correct prediction), this takes a value between 0 and 1
- When  $\exists \mathbf{y} \neq \mathbf{y}^m$  such that  $\mathbf{w} \cdot \mathbf{f}(\mathbf{x}^m, \mathbf{y}) \geq \mathbf{w} \cdot \mathbf{f}(\mathbf{x}^m, \mathbf{y}^m)$  (i.e., incorrect prediction), this takes a value  $\geq 1$
- Thus, this always *upper bounds* the 0-1 loss!
- Minimizing hinge loss is good because it minimizes an upper bound on the 0-1 loss (prediction error)

- It doesn't always make sense to penalize all incorrect predictions equally!
- We can change the constraints to

$$\mathbf{w} \cdot \left( \mathbf{f}(\mathbf{x}^m, \mathbf{y}^m) - \mathbf{f}(\mathbf{x}^m, \mathbf{y}) \right) \geq \Delta(\mathbf{y}, \mathbf{y}^m) - \xi_m, \quad \forall \mathbf{y},$$

where  $\Delta(\mathbf{y}, \mathbf{y}^m) \geq 0$  is a measure of how far the assignment  $\mathbf{y}$  is from the true assignment  $\mathbf{y}^m$

- This is called **margin scaling** (as opposed to *slack scaling*)
- We assume that  $\Delta(\mathbf{y}, \mathbf{y}) = 0$ , which allows us to say that the constraint holds for *all*  $\mathbf{y}$ , rather than just  $\mathbf{y} \neq \mathbf{y}^m$
- A frequently used metric for MRFs is **Hamming distance**, where  $\Delta(\mathbf{y}, \mathbf{y}^m) = \sum_{i \in V} \mathbb{1}[y_i \neq y_i^m]$

# Structural SVM with margin scaling

$$\min_{\mathbf{w}} \sum_m \max_{\mathbf{y} \in \mathcal{Y}} \left( \Delta(\mathbf{y}, \mathbf{y}^m) - \mathbf{w} \cdot \left( \mathbf{f}(\mathbf{x}^m, \mathbf{y}^m) - \mathbf{f}(\mathbf{x}^m, \mathbf{y}) \right) \right) + C \|\mathbf{w}\|^2$$

How to solve this? Many methods!

- 1 Cutting-plane algorithm (Tsochantaridis et al., 2005)
- 2 Stochastic subgradient method (Ratliff et al., 2007)
- 3 Dual Loss Primal Weights algorithm (Meshi et al., 2010)
- 4 Frank-Wolfe algorithm (Lacoste-Julien et al., 2013)



# Stochastic subgradient method

$$\min_{\mathbf{w}} \sum_m \max_{\mathbf{y} \in \mathcal{Y}} \left( \Delta(\mathbf{y}, \mathbf{y}^m) - \mathbf{w} \cdot \left( \mathbf{f}(\mathbf{x}^m, \mathbf{y}^m) - \mathbf{f}(\mathbf{x}^m, \mathbf{y}) \right) \right) + C \|\mathbf{w}\|^2$$

- Although this objective is convex, it is not differentiable everywhere
- We can use a *subgradient* method to minimize (instead of gradient descent)
- The *subgradient* of  $\max_{\mathbf{y} \in \mathcal{Y}} \Delta(\mathbf{y}, \mathbf{y}^m) - \mathbf{w} \cdot \left( \mathbf{f}(\mathbf{x}^m, \mathbf{y}^m) - \mathbf{f}(\mathbf{x}^m, \mathbf{y}) \right)$  at  $\mathbf{w}^{(t)}$  is

$$\mathbf{f}(\mathbf{x}^m, \hat{\mathbf{y}}) - \mathbf{f}(\mathbf{x}^m, \mathbf{y}^m),$$

where  $\hat{\mathbf{y}}$  is one of the maximizers with respect to  $\mathbf{w}^{(t)}$ , i.e.

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y} \in \mathcal{Y}} \Delta(\mathbf{y}, \mathbf{y}^m) + \mathbf{w}^{(t)} \cdot \mathbf{f}(\mathbf{x}^m, \mathbf{y})$$

- This maximization is called *loss-augmented* MAP inference

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y} \in \mathcal{Y}} \Delta(\mathbf{y}, \mathbf{y}^m) + \mathbf{w}^{(t)} \cdot \mathbf{f}(\mathbf{x}^m, \mathbf{y})$$

- When  $\Delta(\mathbf{y}, \mathbf{y}^m) = \sum_{i \in V} \mathbb{1}[y_i \neq y_i^m]$ , this corresponds to adding additional single-node potentials

$$\theta_i(y_i) = 1 \text{ if } y_i \neq y_i^m, \text{ and } 0 \text{ otherwise}$$

- If MAP inference was previously exactly solvable by a combinatorial algorithm, loss-augmented MAP inference typically is too
- The Hamming distance pushes the MAP solution away from the true assignment  $\mathbf{y}^m$

# Cutting-plane algorithm

$$\min_{\mathbf{w}, \xi} \sum_m \xi_m + C \|\mathbf{w}\|^2$$

subject to:

$$\begin{aligned} \mathbf{w} \cdot \left( \mathbf{f}(\mathbf{x}^m, \mathbf{y}^m) - \mathbf{f}(\mathbf{x}^m, \mathbf{y}) \right) &\geq \Delta(\mathbf{y}, \mathbf{y}^m) - \xi_m, \quad \forall m, \mathbf{y} \in \mathcal{Y}_m \\ \xi_m &\geq 0, \quad \forall m \end{aligned}$$

- Start with  $\mathcal{Y}_m = \{\mathbf{y}^m\}$ . Solve for the optimal  $\mathbf{w}^*, \xi^*$
- Then, look to see if any of the *unused* constraints are violated
- To find a violated constraint for data point  $m$ , simply solve the loss-augmented inference problem:

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y} \in \mathcal{Y}} \Delta(\mathbf{y}, \mathbf{y}^m) + \mathbf{w} \cdot \mathbf{f}(\mathbf{x}^m, \mathbf{y})$$

- If  $\hat{\mathbf{y}} \in \mathcal{Y}_m$ , do nothing. Otherwise, let  $\mathcal{Y}_m = \mathcal{Y}_m \cup \{\hat{\mathbf{y}}\}$
- Repeat until no new constraints are added. Then we are optimal!

# Cutting-plane algorithm


- Can prove that, in order to solve the structural SVM up to  $\epsilon$  (additive) accuracy, takes a polynomial number of iterations
- In practice, terminates very quickly

# Summary of convergence rates

Optimization algorithm	Online	Primal/Dual	Type of guarantee	Oracle type	# Oracle calls
dual extragradient (Taskar et al., 2006)	no	primal-‘dual’	saddle point gap	Bregman projection	$O\left(\frac{nR \log  \mathcal{Y} }{\lambda \varepsilon}\right)$
online exponentiated gradient (Collins et al., 2008)	yes	dual	expected dual error	expectation	$O\left(\frac{(n + \log  \mathcal{Y} ) R^2}{\lambda \varepsilon}\right)$
excessive gap reduction (Zhang et al., 2011)	no	primal-dual	duality gap	expectation	$O\left(nR \sqrt{\frac{\log  \mathcal{Y} }{\lambda \varepsilon}}\right)$
BMRM (Teo et al., 2010)	no	primal	$\geq$ primal error	maximization	$O\left(\frac{nR^2}{\lambda \varepsilon}\right)$
1-slack SVM-Struct (Joachims et al., 2009)	no	primal-dual	duality gap	maximization	$O\left(\frac{nR^2}{\lambda \varepsilon}\right)$
stochastic subgradient (Shalev-Shwartz et al., 2010a)	yes	primal	primal error w.h.p.	maximization	$\tilde{O}\left(\frac{R^2}{\lambda \varepsilon}\right)$
this paper: block-coordinate Frank-Wolfe	yes	primal-dual	expected duality gap	maximization	$O\left(\frac{R^2}{\lambda \varepsilon}\right)$ Thm. 3

$R$  same as before.  $n$ =number of training examples.  $\lambda$  is the regularization constant (corresponding to  $2C/n$ )


# Application to segmentation & support inference




Input RGB




Surface Normals



Aligned Normals



Segmentation




Input Depth



Inpainted Depth



3D Planes




Support Relations

(Silberman, Sontag, Fergus. ECCV '14)

# Application to machine translation

Word alignment between languages:



(Taskar, Lacoste-Julien, Klein. EMNLP '05)

# MAP inference

- Recall the MAP inference task,

$$\arg \max_{\mathbf{x}} p(\mathbf{x}), \quad p(\mathbf{x}) = \frac{1}{Z} \prod_{c \in C} \phi_c(\mathbf{x}_c)$$

(we assume any evidence has been subsumed into the potentials)

- Since the normalization term is simply a constant, this is equivalent to

$$\arg \max_{\mathbf{x}} \prod_{c \in C} \phi_c(\mathbf{x}_c)$$

(called the *max-product* inference task)

- Furthermore, since log is monotonic, letting  $\theta_c(\mathbf{x}_c) = \lg \phi_c(\mathbf{x}_c)$ , we have that this is equivalent to


$$\arg \max_{\mathbf{x}} \sum_{c \in C} \theta_c(\mathbf{x}_c)$$

(called *max-sum*)




# Motivating application: image denoising

- Input (left): noisy image
- Output (right): denoised image



# Motivating application: protein side-chain placement


- Find “minimum energy” conformation of amino acid side-chains along a fixed carbon backbone:



- Orientations of the side-chains are represented by discretized angles called rotamers
- Rotamer choices for nearby amino acids are energetically coupled (attractive and repulsive forces)

# Motivating application: dependency parsing

- Given a sentence, predict the dependency tree that relates the words:



- Arc from head word of each phrase to words that modify it
- May be *non-projective*: each word and its descendants may not be a contiguous subsequence
- $m$  words  $\implies m(m-1)$  binary arc selection variables  $x_{ij} \in \{0, 1\}$
- Let  $\mathbf{x}_{|i} = \{x_{ij}\}_{j \neq i}$  (all outgoing edges). Predict with:

$$\max_{\mathbf{x}} \theta_T(\mathbf{x}) + \sum_{ij} \theta_{ij}(x_{ij}) + \sum_i \theta_{i|}(\mathbf{x}_{|i})$$

# How to perform approximate MAP inference?

- Local search (*iterated conditional modes*)
  - Start from an arbitrary assignment (e.g., random). Iterate:
  - Choose a variable. Change a new state for this variable to maximize the value of the resulting assignment
- Branch-and-bound
  - Exhaustive search over space of assignments, pruning branches that can be provably shown not to contain a MAP assignment
  - Can use LP relaxations or dual decomposition to obtain upper bounds
  - Lower bound obtained from value of any assignment found
- Branch-and-cut (most powerful method; used by CPLEX & Gurobi)
  - Same as branch-and-bound; spend more time getting tighter bounds
  - Adds *new constraints* to cut off fractional solutions of the LP relaxation, making the upper bound tighter

# Dual decomposition

*effective, simple, approximate MAP inference, with guarantees*

- Consider the MAP problem for pairwise Markov random fields:

$$\text{MAP}(\theta) = \max_{\mathbf{x}} \sum_{i \in V} \theta_i(x_i) + \sum_{ij \in E} \theta_{ij}(x_i, x_j).$$

- If we push the maximizations *inside* the sums, the value can only *increase*:



$$\text{MAP}(\theta) \leq \sum_{i \in V} \max_{x_i} \theta_i(x_i) + \sum_{ij \in E} \max_{x_i, x_j} \theta_{ij}(x_i, x_j)$$

- Note that the right-hand side can be easily evaluated
- One can always *reparameterize* a distribution by operations like

$$\begin{aligned}\theta_i^{\text{new}}(x_i) &= \theta_i^{\text{old}}(x_i) + f(x_i) \\ \theta_{ij}^{\text{new}}(x_i, x_j) &= \theta_{ij}^{\text{old}}(x_i, x_j) - f(x_i)\end{aligned}$$

for **any** function  $f(x_i)$ , without changing the distribution/energy

# Dual decomposition



# Dual decomposition

- Define:

$$\tilde{\theta}_i(x_i) = \theta_i(x_i) + \sum_{ij \in E} \delta_{j \rightarrow i}(x_i)$$

$$\tilde{\theta}_{ij}(x_i, x_j) = \theta_{ij}(x_i, x_j) - \delta_{j \rightarrow i}(x_i) - \delta_{i \rightarrow j}(x_j)$$

- It is easy to verify that

$$\sum_i \theta_i(x_i) + \sum_{ij \in E} \theta_{ij}(x_i, x_j) = \sum_i \tilde{\theta}_i(x_i) + \sum_{ij \in E} \tilde{\theta}_{ij}(x_i, x_j) \quad \forall \mathbf{x}$$

- Thus, we have that:

$$\text{MAP}(\theta) = \text{MAP}(\tilde{\theta}) \leq \sum_{i \in V} \max_{x_i} \tilde{\theta}_i(x_i) + \sum_{ij \in E} \max_{x_i, x_j} \tilde{\theta}_{ij}(x_i, x_j)$$

- Every value of  $\delta$  gives a different upper bound on the value of the MAP!
- The **tightest** upper bound can be obtained by minimizing the r.h.s. with respect to  $\delta$ !

# Dual decomposition

- We obtain the following **dual** objective:  $L(\delta) =$

$$\sum_{i \in V} \max_{x_i} \left( \theta_i(x_i) + \sum_{ij \in E} \delta_{j \rightarrow i}(x_i) \right) + \sum_{ij \in E} \max_{x_i, x_j} \left( \theta_{ij}(x_i, x_j) - \delta_{j \rightarrow i}(x_i) - \delta_{i \rightarrow j}(x_j) \right),$$

$$\text{DUAL-LP}(\theta) = \min_{\delta} L(\delta)$$

- This provides an upper bound on the MAP assignment!

$$\text{MAP}(\theta) \leq \text{DUAL-LP}(\theta) \leq L(\delta)$$

- How can find  $\delta$  which give tight bounds?




# Solving the dual efficiently

- Many ways to solve the dual linear program, i.e. minimize with respect to  $\delta$ :

$$\sum_{i \in V} \max_{x_i} \left( \theta_i(x_i) + \sum_{ij \in E} \delta_{j \rightarrow i}(x_i) \right) + \sum_{ij \in E} \max_{x_i, x_j} \left( \theta_{ij}(x_i, x_j) - \delta_{j \rightarrow i}(x_i) - \delta_{i \rightarrow j}(x_j) \right),$$

- One option is to use the subgradient method
- Can also solve using **block coordinate-descent**, which gives algorithms that look very much like belief propagation:



# Max-product linear programming (MPLP) algorithm

**Input:** A set of factors  $\theta_i(x_i), \theta_{ij}(x_i, x_j)$

**Output:** An assignment  $x_1, \dots, x_n$  that approximates the MAP

**Algorithm:**

- Initialize  $\delta_{i \rightarrow j}(x_j) = 0, \quad \delta_{j \rightarrow i}(x_i) = 0, \quad \forall ij \in E, x_i, x_j$
- Iterate until small enough change in  $L(\delta)$ :  
For each edge  $ij \in E$  (sequentially), perform the updates:

$$\delta_{j \rightarrow i}(x_i) = -\frac{1}{2}\delta_i^{-j}(x_i) + \frac{1}{2} \max_{x_j} \left[ \theta_{ij}(x_i, x_j) + \delta_j^{-i}(x_j) \right] \quad \forall x_i$$

$$\delta_{i \rightarrow j}(x_j) = -\frac{1}{2}\delta_j^{-i}(x_j) + \frac{1}{2} \max_{x_i} \left[ \theta_{ij}(x_i, x_j) + \delta_i^{-j}(x_i) \right] \quad \forall x_j$$

where  $\delta_i^{-j}(x_i) = \theta_i(x_i) + \sum_{ik \in E, k \neq j} \delta_{k \rightarrow i}(x_i)$

- Return  $x_i \in \arg \max_{\hat{x}_i} \tilde{\theta}_i^{\delta}(\hat{x}_i)$

# Generalization to arbitrary factor graphs

## Inputs:

- A set of factors  $\theta_i(x_i), \theta_f(\mathbf{x}_f)$ .

## Output:

- An assignment  $x_1, \dots, x_n$  that approximates the MAP.

## Algorithm:

- Initialize  $\delta_{fi}(x_i) = 0, \quad \forall f \in F, i \in f, x_i$ .
- Iterate until small enough change in  $L(\boldsymbol{\delta})$  (see Eq. 1.2):  
For each  $f \in F$ , perform the updates

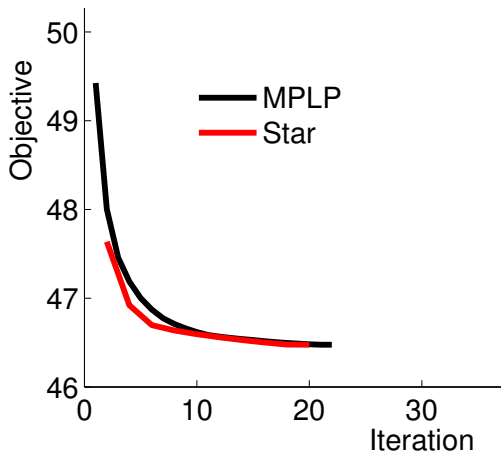
$$\delta_{fi}(x_i) = -\delta_i^{-f}(x_i) + \frac{1}{|f|} \max_{\mathbf{x}_{f \setminus i}} \left[ \theta_f(\mathbf{x}_f) + \sum_{\hat{i} \in f} \delta_{\hat{i}}^{-f}(x_{\hat{i}}) \right], \quad (1.16)$$

simultaneously for all  $i \in f$  and  $x_i$ . We define  $\delta_i^{-f}(x_i) = \theta_i(x_i) + \sum_{f \neq f} \delta_{fi}(x_i)$ .

- Return  $x_i \in \arg \max_{\hat{x}_i} \bar{\theta}_i^{\boldsymbol{\delta}}(\hat{x}_i)$  (see Eq. 1.6).


# Experimental results

Comparison of two block coordinate descent algorithms on a  $10 \times 10$  node Ising grid:




# Experimental results

Performance on stereo vision inference task:



# Linear programming duality

Beyond the scope of this class, but one can show intimate relationship between dual decomposition and linear programming relaxations:



$$\text{MAP}(\theta) \leq \text{LP}(\theta) = \text{DUAL-LP}(\theta) \leq L(\delta)$$